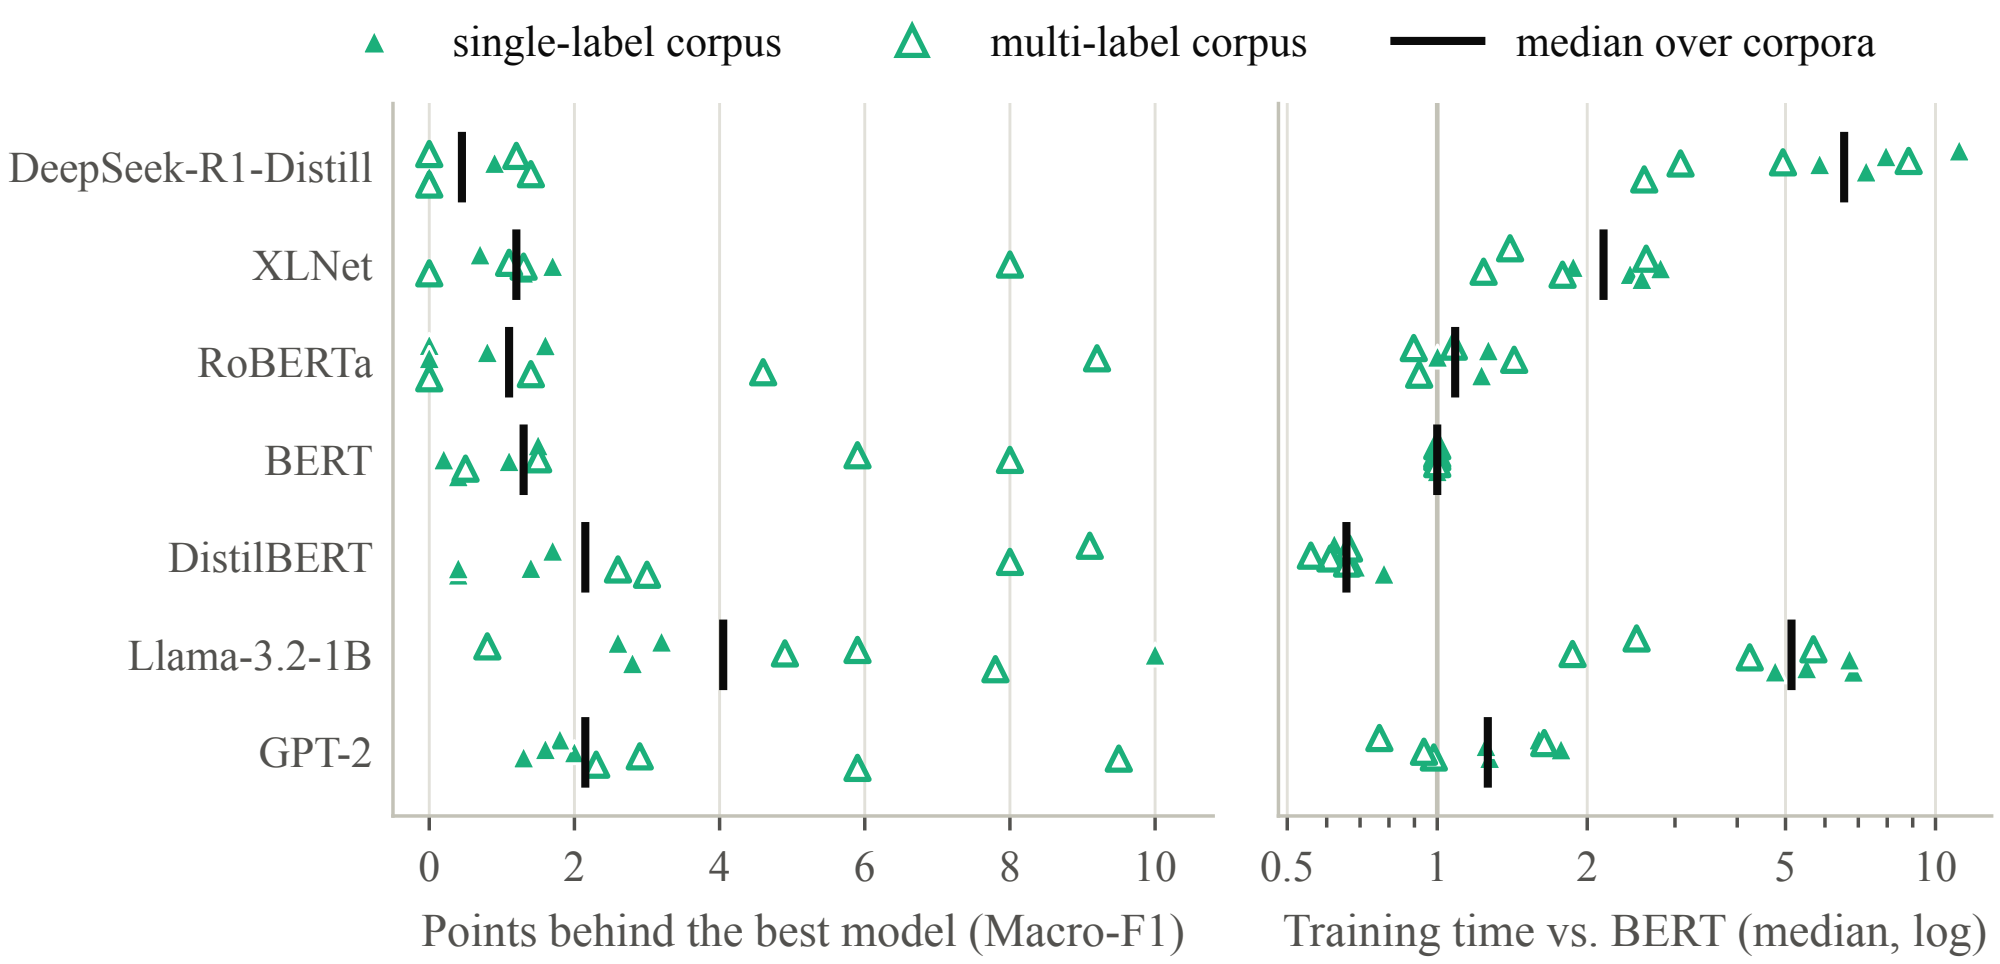
single-label corpus
multi-label corpus
median over corpora
DeepSeek-R1-Distill
XLNet
RoBERTa
BERT
DistilBERT
Llama-3.2-1B
GPT-2
0
2
4
6
8
10
Points behind the best model (Macro-F1)
0.5
1
2
5
10
Training time vs. BERT (median, log)